



Logistic Regression python (**scikitlearn**)

MACHINE LEARNING

사이킷런을 사용해 로지스틱 회귀 모델 훈련하기



```
from sklearn.linear_model import LogisticRegression
```

```
lr = LogisticRegression(solver='liblinear', multi_class='auto', C=100.0, random_state=1)  
lr.fit(X_train_stand, y_train)
```

```
X_combined_stand = np.vstack((X_train_stand, X_test_stand))  
y_combined = np.hstack((y_train, y_test))
```

```
plot_decision_regions(X_combined_stand, y_combined,  
                      classifier=lr, test_idx=range(105, 150))
```

```
plt.xlabel('petal length [standardized]')  
plt.ylabel('petal width [standardized]')  
plt.legend(loc='upper left')  
plt.tight_layout()  
plt.show()
```

`sklearn.linear_model`: Linear Models

The `sklearn.linear_model` module implements a variety of linear models.

User guide: See the [Linear Models](#) section for further details.

The following subsections are only rough guidelines: the same estimator can fall into multiple categories, depending on its parameters.

Linear classifiers

| | |
|--|--|
| <code>linear_model.LogisticRegression([penalty, ...])</code> | Logistic Regression (aka logit, MaxEnt) classifier. |
| <code>linear_model.LogisticRegressionCV(*[, Cs, ...])</code> | Logistic Regression CV (aka logit, MaxEnt) classifier. |
| <code>linear_model.PassiveAggressiveClassifier(*)</code> | Passive Aggressive Classifier |
| <code>linear_model.Perceptron(*[, penalty, alpha, ...])</code> | Read more in the User Guide . |
| <code>linear_model.RidgeClassifier([alpha, ...])</code> | Classifier using Ridge regression. |
| <code>linear_model.RidgeClassifierCV([alphas, ...])</code> | Ridge classifier with built-in cross-validation. |
| <code>linear_model.SGDClassifier([loss, penalty, ...])</code> | Linear classifiers (SVM, logistic regression, etc.) with SGD training. |

Classical linear regressors

| | |
|--|---|
| <code>linear_model.LinearRegression(*[, ...])</code> | Ordinary least squares Linear Regression. |
| <code>linear_model.Ridge([alpha, fit_intercept, ...])</code> | Linear least squares with l2 regularization. |
| <code>linear_model.RidgeCV([alphas, ...])</code> | Ridge regression with built-in cross-validation. |
| <code>linear_model.SGDRegressor([loss, penalty, ...])</code> | Linear model fitted by minimizing a regularized empirical loss with SGD |

Solver

최적화의 문제를 풀어내는 이론들

| | Solvers | | | | |
|------------------------------|-------------|---------|-------------|-------|--------|
| Penalties | 'liblinear' | 'lbfgs' | 'newton-cg' | 'sag' | 'saga' |
| Multinomial + L2 penalty | no | yes | yes | yes | yes |
| OVR + L2 penalty | yes | yes | yes | yes | yes |
| Multinomial + L1 penalty | no | no | no | no | yes |
| OVR + L1 penalty | yes | no | no | no | yes |
| Elastic-Net | no | no | no | no | yes |
| No penalty ('none') | no | yes | yes | yes | yes |
| Behaviors | | | | | |
| Penalize the intercept (bad) | yes | no | no | no | no |
| Faster for large datasets | no | no | no | yes | yes |
| Robust to unscaled datasets | yes | yes | yes | no | no |

Solver

Limited-memory Broyden-Fletcher-Goldfarb-Shanno
Algorithm

Multinomial 다항의

OVR = One vs Rest

OVO = One vs One

topmost

모두다 약자 (두문자어, acronym)

| | Solvers | | | | |
|------------------------------|-------------|---------|-------------|-------|--------|
| Penalties | 'liblinear' | 'lbfgs' | 'newton-cg' | 'sag' | 'saga' |
| Multinomial + L2 penalty | no | yes | yes | yes | yes |
| OVR + L2 penalty | yes | yes | yes | yes | yes |
| Multinomial + L1 penalty | no | no | no | no | yes |
| OVR + L1 penalty | yes | no | no | no | yes |
| Elastic-Net | no | no | no | no | yes |
| No penalty ('none') | no | yes | yes | yes | yes |
| Behaviors | | | | | |
| Penalize the intercept (bad) | yes | no | no | no | no |
| Faster for large datasets | no | no | no | yes | yes |
| Robust to unscaled datasets | yes | yes | yes | no | no |

Solver

The “saga” solver is often the best choice.

The “liblinear” solver is used by default for historical reasons.

LIBLINEAR – the winner of ICML 2008 large-scale learning challenge. It applies Automatic parameter selection

‘A **Lib**rary for **Large Linear** Classification’ <https://www.csie.ntu.edu.tw/~cjlin/liblinear/>

SAG – Mark Schmidt, Nicolas Le Roux, and Francis Bach (2013).

‘Minimizing Finite Sums with the **S**tochastic **A**verage **G**radient’ <https://hal.inria.fr/hal-00860051/document>

SAGA – Defazio, A., Bach F. & Lacoste-Julien S. (2014). (a variant of SAG that also supports L1 option (**L2** or **L1**))

‘SAGA: A Fast Incremental Gradient Method With Support for Non-Strongly Convex Composite Objectives’ <https://arxiv.org/abs/1407.0202>

Solver

최적화의 문제를 풀어내는 이론들 → 결국 데이터에 맞는 방법 찾기

OVR = One vs Rest

| | Solvers | | | | |
|------------------------------|-------------|---------|-------------|-------|--------|
| Penalties | 'liblinear' | 'lbfgs' | 'newton-cg' | 'sag' | 'saga' |
| Multinomial + L2 penalty ← | no | yes | yes | yes | yes |
| OVR + L2 penalty ← | yes | yes | yes | yes | yes |
| Multinomial + L1 penalty ← | no | no | no | no | yes |
| OVR + L1 penalty ← | yes | no | no | no | yes |
| Elastic-Net | no | no | no | no | yes |
| No penalty ('none') | no | yes | yes | yes | yes |
| Behaviors | | | | | |
| Penalize the intercept (bad) | yes | no | no | no | no |
| Faster for large datasets | no | no | no | yes | yes |
| Robust to unscaled datasets | yes | yes | yes | no | no |

pipeline module. Pipeline class

```
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
```

`sklearn.pipeline.Pipeline`

`class sklearn.pipeline.Pipeline(steps, *, memory=None, verbose=False)` Verb 동사 말 많은,
= 지금 상황 다 설명해줘

데이터 처리, 분류 등의 steps단계를 묶어서 라인으로 만든 class의
인스턴스를 만들 수 있다.

```
lr_tfidf = Pipeline([('vect', tfidf),  
                     ('clf', LogisticRegression(solver='liblinear', random_state=0))])
```




feature_extraction module

```
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.feature_extraction.text import TfidfVectorizer
```

TfidfVectorizer

: 문서를 tf-idf의 **feature matrix**로 벡터변환하는 클래스

BOW (Bag of Words)

고정된 bag (multiset) 가방 자리 를 만들고

D_i 라는 개별 문서의 가방 자리에 해당하는 단어들이 포함되어 있는지 표시

I love this movie! It's sweet,
but with satirical humor. The
dialogue is great and the
adventure scenes are fun...
It manages to be whimsical
and romantic while laughing
at the conventions of the
fairy tale genre. I would
recommend it to just about
anyone. I've seen it several
times, and I'm always happy
to see it again whenever I
have a friend who hasn't
seen it yet!



| | |
|-----------|-----|
| it | 6 |
| I | 5 |
| the | 4 |
| to | 3 |
| and | 3 |
| seen | 2 |
| yet | 1 |
| would | 1 |
| whimsical | 1 |
| times | 1 |
| sweet | 1 |
| satirical | 1 |
| adventure | 1 |
| genre | 1 |
| fairy | 1 |
| humor | 1 |
| have | 1 |
| great | 1 |
| ... | ... |

Vectorizer 벡터화 class

CountVectorizer:

문서 집합에서 단어 토큰을 생성하고
각 단어의 수를 세어 BOW 인코딩한 벡터를 만든다.

TfidfVectorizer:

TF-IDF 방식으로 단어의 중요도를 조정한 BOW 벡터를 만든다.

```
sklearn.feature_extraction.text.TfidfVectorizer
```

```
class sklearn.feature_extraction.text.TfidfVectorizer(*, input='content', encoding='utf-8', decode_error='strict',
```

`sklearn.feature_extraction.text` submodule gathers utilities to build feature vectors from text documents.

`sklearn.feature_extraction.text.TfidfVectorizer`

```
class sklearn.feature_extraction.text.TfidfVectorizer(*, input='content', encoding='utf-8', decode_error='strict',
```

TfidfVectorizer:

TF-IDF 방식으로 단어의 가중치를 조정한 BOW 벡터를 만든다.

단순 단어 빈도로 접근하는 게 아니라,

어떤 단어가 한 문서에서 많이 나타난 동시에

다른 문서에서는 잘 나타나지 않는 것까지 고려하기 위한 개념

TF-IDF (Term Frequency-Inverse Document Frequency)